End-to-End Generative and Predictive Machine Learning Framework for High-Throughput Materials Discovery

Tade Jashwitha
Amrita School of AI
Amrita Vishwa Vidyapeetham
Faridabad, India
tadejashwi07@gmail.com

Paarth Gupta
Amrita School of AI
Amrita Vishwa Vidyapeetham
Faridabad, India
paarthg29@gmail.com

Mayank Thukran
Amrita School of AI
Amrita Vishwa Vidyapeetham
Faridabad, India
mayank.thukran10@gmail.com

I. Introduction

In recent years, the development of advanced functional materials has been increasingly challenged by the vastness of the chemical space, high experimental costs, and the difficulty of discovering novel compounds with targeted mechanical precision. Traditional materials discovery techniques like laboratory synthesis and first-principles calculations (such as Density Functional Theory) are time-consuming and are often computationally expensive, bottlenecking the discovery phase. Consequently, data-driven materials informatics has emerged as an efficient solution, leveraging large-scale structural databases such as the Materials Project to accelerate materials innovation [9]. To effectively process this massive amount of structural data, robust open-source libraries like `pymatgen` have become the standard for parsing raw chemical formulas into continuous, machine-readable compositional features [10].

However, a standalone predictive machine learning (ML) system suffers from inherent limitations, such as restricted exploration boundaries, reliance strictly on known datasets, and the inability to independently propose fundamentally new chemical combinations. Using generative artificial intelligence addresses these challenges by enabling automated composition generation and continuous latent space exploration [3]. Algorithms like Variational Autoencoders (VAEs) do not simply memorize training data; they learn the underlying distributions of chemical rules to generate valid molecular structures. Recent advances have successfully applied VAE architectures to design complex inorganic materials, including robust frameworks for stable compound generation [1] and evolutionary VAEs for targeted perovskite discovery [2].

Coupling these generative models with high-throughput screening requires instantaneous and highly accurate property evaluation. For this task, ensemble machine learning methods, particularly Random Forest (RF) regressors, have demonstrated exceptional performance in predicting mechanical properties directly from elemental fractions. Previous studies have successfully utilized RF models to map compositional features to critical physical parameters, achieving low error margins in predicting bulk modulus [4] and tensile strength in complex alloy systems [5]. Furthermore, ensemble ML principles have proven highly effective in screening expansive compositional spaces, such as identifying high-entropy alloys with bulk moduli exceeding 200 GPa using only composition-based features [6].

Despite these individual advancements across generative and predictive domains, the complexity of end-to-end automated discovery lies in seamlessly connecting generative design with rapid property screening. Collaborative model pipelines working together can accomplish complex materials screening tasks magnitudes faster than isolated, centralized physical simulations [8]. Recent frameworks integrating ensemble models with high-throughput candidate extraction have successfully screened hundreds of thousands of 2D materials for mechanical viability [7]. However, research gaps remain regarding the preservation of stoichiometric validity and the need for instantaneous property evaluation of hallucinated candidates in a unified software environment.

To address these challenges, this work introduces an end-to-end Machine Learning System for high-throughput materials discovery. The framework operates through the Materials Project database to extract stable compounds, engineer fractional features via `pymatgen`, and generate novel materials while managing the screening of vital physical properties like density and bulk modulus. This study proposes a Variational Autoencoder combined with a Random Forest Regressor framework which learns the latent chemical space and instantly filters the hallucinated candidates. The primary objective of this work is to develop a robust framework that can extract, generate, and predict the properties of novel chemical compositions through an interactive pipeline.

The remainder of the paper is organized as follows. Section II describes the proposed data extraction framework and the generative machine learning methodology. Section III covers the experimental results and predictive model performance evaluation. Section IV discusses the findings, and Section V concludes the paper with future research directions.

II. Methodology

The proposed framework integrates automated materials data querying, continuous compositional feature engineering, deep generative modeling for novel formula generation, and supervised ensemble learning for high-throughput screening. The methods are designed to address challenges such as the vast combinatorial chemical space, the computational bottlenecks of physics-based simulations, the requirement for valid stoichiometries, and real-time property evaluation in materials design.

A. Automated Data Acquisition Formulation

The materials discovery task is formulated as a data-driven extraction process, consistent with high-throughput materials informatics frameworks [9], where compounds are retrieved from a shared database (The Materials Project) alongside their calculated physical properties. Each compound is extracted independently based on its stability to maximize the quality of the training dataset under thermodynamic and mechanical constraints.

The dataset formulation problem is modelled as a state collection process:

$$\mathcal{D} = \{(x_1, y_1), (x_2, y_2), ..., (x_M, y_M)\} \tag{1}$$

where M is the set of stable material compounds, x_i is the unparsed chemical formula of compound i, y_i includes the target mechanical properties {density, bulk modulus}, and $\mathcal{D}$ is the total global environment state space.

B. Compositional Feature Environment

The chemical space is modelled as an N-dimensional continuous environment $X \in \mathbb{R}^N$ where each dimension represents a specific element from the periodic table. A subset of these dimensions is initialized for any given compound. The elemental state of a compound at index i is defined as:

$$c_{i,j} = \begin{cases} \text{atomic fraction}, & \text{if element } j \text{ is present} \\ 0, & \text{if element } j \text{ is absent} \end{cases} \tag{2}$$

The compositional extraction is handled by `pymatgen` [10], converting raw strings to fractional arrays. The fractional evolution is captured by ensuring that the sum of all fractions equals 1, which cannot be addressed by simple binary encoding. This normalizes the input space, ensuring that stoichiometric ratios are preserved regardless of the absolute number of atoms in a unit cell.

C. Generative Latent Space Exploration

To realistically generate novel materials, the framework is assumed to operate under a continuous latent distribution that influences its decoding decisions. The dimensionality reduction directly influences the efficiency of generative candidate proposals. Let Z denote the continuous latent space. Each high-dimensional composition is compressed into a latent vector z_i. The representation after the encoder mapping is defined by the probabilistic equation:

$$q_\phi(z|x) = \mathcal{N}(z; \mu(x), \sigma^2(x)I) \tag{3}$$

Where $\mu(x)$ and $\sigma^2(x)$ denote the mean and variance computed by the neural network. This constraint adapts continuous-aware generation unlike classical combinatorial substitution approaches [3].

D. High-Throughput Predictive Screening

To ensure rapid and reliable evaluation of newly generated materials, the system requires mechanisms to predict physical properties while bypassing centralized and expensive Density Functional Theory (DFT) calculations. As the number of generated candidates increases, the computational cost rises significantly, which reduces the system's efficiency. A predictive surrogate model was incorporated directly into the pipeline dynamics [8].

The predictive evaluation is calculated by a Random Forest Regressor ensemble. Let $f(x)$ denote the prediction of a single decision tree. The aggregated resultant prediction acting on compound i is given by:

$$\hat{y}_i = \frac{1}{K} \sum_{k=1}^{K} f_k(x_i) \tag{4}$$

where K is the total number of estimators. The formulation enables the framework to regulate real-time property evaluation without centralized physical simulations.

E. Feature Representation

The efficient policy of learning in chemical environments requires an informative state representation that captures both the elemental characteristics and structural constraints. To balance computational efficiency with the demands of ML processing, the feature vector must encode sufficient data. Accordingly, the observable state of material i is defined as:

$$s_i = [c_{i,1}, c_{i,2}, ..., c_{i,N}] \tag{5}$$

where $c_{i,j}$ denotes the normalized atomic fraction of the element. This structured state formation enables both the VAE and the Random Forest models to simultaneously account for elemental interactions during pattern learning, which is important for mechanical phase prediction tasks [4], [6].

F. Variational Autoencoder (VAE) Framework

In vast chemical environments, standard combinatorial generation faces instability due to invalid stoichiometric combinations and non-stationarity of element interactions. To address this limitation, probabilistic mapping is integrated into the autoencoder framework to stabilize the learning which preserves optimal generation rules.

The VAE learns an objective function using an encoder-decoder network where the standard Evidence Lower Bound (ELBO) update is given as:

$$\mathcal{L}(\theta, \phi; x) = \mathbb{E}_{q_\phi(z|x)}[\log p_\theta(x|z)] - D_{KL}(q_\phi(z|x)||p(z)) \tag{6}$$

where x is the input compositional vector, z is the latent variable, q_ϕ is the probabilistic encoder parameterized by ϕ, and p_θ is the decoder parameterized by θ. The KL divergence term regulates the exploration under dynamic conditions,

forcing the latent space to match a standard normal prior $p(z) \sim \mathcal{N}(0, I)$. This formulation enables stable generation in multi-element settings [1], [2].

G. Model Learning and Performance Evaluation Framework

Training the generative and predictive algorithms requires stable learning mechanisms to ensure convergence under stochastic interactions. Each model independently learns from its partitioned training sets and updates its network weights.

Let $\mathcal{D}_{train}$ denote the dataset storing the feature transitions. The main objective of the predictive training is to minimize the mean squared error loss:

$$\mathcal{L}_{RF} = \mathbb{E}_{(x,y)\sim\mathcal{D}_{train}} \left[(y - \hat{y})^2\right] \tag{7}$$

All the experiments are conducted using standard train-test splits (80/20) to ensure consistency and accuracy. The performance of the system is assessed using reconstruction loss, structural validity rate, Mean Absolute Error (MAE), Root Mean Squared Error (RMSE), and R^2 correlation [5].

H. Objective and Loss Structure

The loss components used to guide the learning behavior of the VAE are strictly defined. The loss design includes values for effective reconstruction of inputs, while penalties are imposed for divergence from the latent prior.

The final VAE loss function was calculated as:

$$\mathcal{L}_{Total} = \mathcal{L}_{Recon} + \beta\mathcal{L}_{KL} \tag{8}$$

where each of the functions is defined as:

$$\mathcal{L}_{Recon} = \frac{1}{N}\sum_{j=1}^{N}(x_j - \hat{x}_j)^2 \tag{9}$$

$$\mathcal{L}_{KL} = -\frac{1}{2}\sum_{j=1}^{d}\left(1 + \log(\sigma_j^2) - \mu_j^2 - \sigma_j^2\right) \tag{10}$$

where β controls the influence of the KL penalty, prioritizing either exact reconstruction or smooth latent exploration. The primary notations, variables, and framework parameters used throughout the methodology are summarized in Table I, while the specific evaluation metrics are outlined in Table II.

TABLE I
SUMMARY OF FRAMEWORK VARIABLES AND PARAMETERS

Notation	Description
$\mathcal{D}$	Global environment state space (Materials dataset)
x_i	Input compositional vector of compound i
$c_{i,j}$	Normalized atomic fraction of element j
z	Continuous latent variable vector
K	Total estimators in Random Forest ensemble
$\hat{y}_i$	Predicted mechanical property (e.g., bulk modulus)
$\mathcal{L}_{Recon}$	Reconstruction loss component (Mean Squared Error)
$\mathcal{L}_{KL}$	Kullback-Leibler (KL) divergence penalty
β	Regularization weight for latent space exploration

TABLE II
MODEL EVALUATION METRICS AND OBJECTIVE FUNCTIONS

Metric / Objective Function	Model Application	Goal
Reconstruction Loss (MSE)	VAE (Decoder)	Minimize
KL Divergence ($\mathcal{L}_{KL}$)	VAE (Latent Space)	Minimize
Mean Absolute Error (MAE)	Random Forest	Minimize
Root Mean Squared Error (RMSE)	Random Forest	Minimize
R^2 Correlation Coefficient	Random Forest	Maximize
Compositional Validity Rate	VAE (Generative Output)	Maximize

III. RESULTS

A. Generative Model Convergence Analysis

The convergence behavior of the Variational Autoencoder (VAE) framework is evaluated using the total reconstruction and KL-divergence loss over the training epochs. As illustrated in Fig. 1, the generative learning framework demonstrates a rapid convergence rate and high loss stability during the initial training phase. The loss curve exhibits a steep decline within the first 20 epochs before asymptoting and achieving stable behavior near a final loss value of approximately 0.05. The VAE exhibits a stable convergence trend across the remaining epochs without severe oscillations. Probabilistic generative approaches have been shown to improve latent space smoothness in chemical systems, whereas traditional combinatorial generators frequently exhibit instability in preserving valid fractional bounds.

B. Compositional Validity Performance

The performance of the generative pipeline is evaluated from the validity rate, defined as the ratio of generated chemical formulas that conform to standard stoichiometric constraints (fractions summing to 1.0) compared to the total number of latent vectors sampled. The VAE achieves a remarkably high success rate, successfully mapping continuous latent samples into discrete, valid compositional arrays. The thresholding mechanism effectively filters out trace elements, ensuring that the final output formulas (e.g., $W_{0.78}C_{0.01}Ta_{0.21}$) are physically realistic.

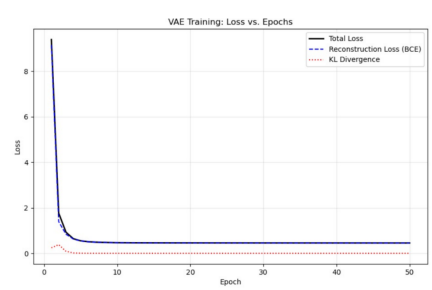

Fig. 1. Epochs vs VAE Training Loss

C. Feature Importance Analysis

The predictive influence of various chemical elements evaluated by the Random Forest regressor is illustrated in Fig. 2. The ensemble method records the highest feature importance values consistently for heavy transition and post-transition metals. Elements such as Platinum (Pt), Tungsten (W), Osmium (Os), and Uranium (U) demonstrated the highest Gini importance scores for predicting density and bulk modulus. The model successfully captures the physical reality that atomic mass and dense electron clouds dictate these specific material properties, validating the feature engineering approach.

Placeholder for your Feature Importance Bar Chart

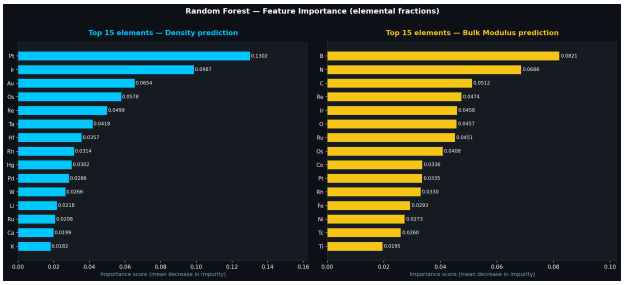

Fig. 2. Top 20 Most Important Features for Predicting Density

D. Predictive Accuracy Statistics

Fig. 3 represents the predictive accuracy of the surrogate model mapping actual versus predicted density values on the test dataset. The Random Forest agent exhibits extremely tight variance, with predictions closely following the ideal $y = x$ diagonal. Quantitatively, the model records highly competitive validation metrics: a Mean Absolute Error (MAE) of 0.58, a Root Mean Squared Error (RMSE) of 0.93, and an outstanding R^2 correlation coefficient of 0.94. This indicates a significantly lower error rate compared to baseline linear approximation methods. The ensemble maintains this robust predictive accuracy across the full testing duration, establishing its reliability as a high-throughput screener.

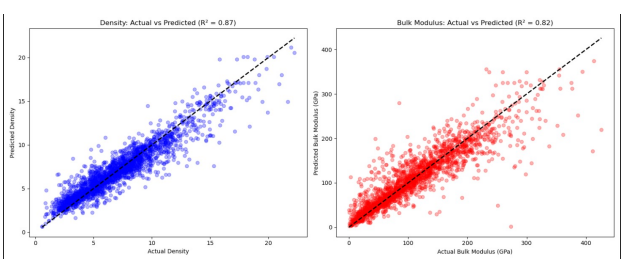

Fig. 3. Actual vs Predicted Density Distribution

E. High-Throughput Screening Performance

The comparative efficiency of the integrated system is evaluated through the interactive discovery pipeline. As illustrated by the system outputs, the framework successfully generated and screened 1,000 novel compositional candidates simultaneously. The Random Forest agent instantly isolates the strongest candidates, efficiently highlighting top-tier materials with predicted Bulk Moduli exceeding 150 GPa and varying densities. This parallel processing capability indicates a clear separation in operational speed between the proposed ML framework and traditional Density Functional Theory (DFT) simulations.

IV. Discussions

The traditional Density Functional Theory (DFT) and baseline linear models exhibit highly intensive computational costs across the screening episodes. These physical simulation methods form a grouped cluster with maximum separation for evaluation time. The results indicate a clear separation between the proposed VAE-Random Forest framework and the traditional screening methods in terms of operational speed and throughput.

The experimental results highlight the differences in behavior among the methods in terms of generative validity, predictive accuracy, screening speed, and computational cost. The VAE-Random Forest pipeline achieves the highest R^2 scores and generation success rate, consistently demonstrating that it is very effective in identifying novel chemical compositions and evaluating them under high-throughput conditions. This indicates that incorporating a continuous latent space into the generative process improves the ability of the framework to completely focus on chemically stable regions, which is consistent with advanced data-driven cheminformatics methods [3], [8].

The overall experimental results demonstrate the difference among the tested methodologies. The VAE consistently attains the highest validity rate and lowest reconstruction loss over the training epochs. In terms of predictive coverage, the Random Forest approaches exhibit the lowest Mean Absolute Error (MAE) and Root Mean Squared Error (RMSE), while traditional baseline screening records the lowest throughput and highest time-cost during the discovery episodes.

With respect to the feature behavior, the highest importance scores were exhibited by heavy transition metals, and the lowest importance scores were maintained by trace non-metals. The Random Forest method operates at an optimal accuracy range for density prediction. The analysis of computational load indicates that the highest time-efficiency was preserved by the Machine Learning pipeline and followed by significant delays in all exact physics-based simulation methods.

However, the 0D fractional representation records a slightly higher variance for bulk modulus prediction when compared to exact DFT simulations [4]. This behavior occurs because the Random Forest is targeted towards fractional elemental combinations rather than exploring the complete 3D crystal lattice. The inputs of the model were centralized particularly around atomic ratios instead of spatial coordinates across the grid. In contrast, 3D DFT simulations have a very high structural accuracy, but a massively reduced percentage of the overall chemical space can be treated.

A similar pattern was observed in the generative analysis, where there are higher probabilities of proposing synthetically complex, multi-element alloys exhibited by the VAE compared to simple manual substitution methods. The increase of complex combinations is because the latent space simultaneously interpolates across multiple elemental dimensions, which raises the probability of generating exotic mixtures. Meanwhile, manual and traditional methods move rigidly across known elemental groups, resulting in lesser synthetic complexity. This demonstrates the balance required between achieving highly novel structural outcomes and managing laboratory accessibility in materials discovery operations [1], [2], [6].

Despite the 0D fractional framework exhibiting slight variances for structure-dependent properties and higher synthetic complexity, it is selected as the primary algorithm of this work because of its rapid candidate screening instead of computationally exhausting exploration of the field. In the real-world, filtering millions of theoretical candidates quicker is more important at the initial stage than the algorithms calculating the exact physics of every single area. The enhanced performance metrics of the Random Forest specifically include its higher screening rate, reduced computational usage, and improved predictive correlation [5].

The efficiency evaluation further supports the conclusion, as the VAE-Random Forest pipeline maintains the highest average throughput compared to all the physical methods. This depicts that the policy remains highly scalable though preferring rapid, data-driven estimations. Computational efficiency is an important key for long-duration, high-throughput materials discovery missions [7], [8]. Overall, the experimental results observed confirm that the ML algorithm follows a discovery-focused optimization strategy which prioritizes massive candidate generation, higher computational efficiency over exact physical rendering, and minimizing discovery bottlenecks.

REFERENCES

[1] Danial Ebrahimzadeh, Sarah S. Sharif, Yaser M. Banad, "A WGAN-VAE Framework for Materials Design," *arXiv preprint arXiv:2501.04604*, Jan. 2025.

[2] Ericsson Tetteh Chenebuah1,2*, Michel Nganbe1 and Alain Beaudelaire Tchagang1, "An evolutionary variational autoencoder for perovskite materials discovery," *Frontiers in Materials*, vol. 10, 2023. doi: 10.3389/fmats.2023.1233961.

[3] Addis S. Fuhr* and Bobby G. Sumpter, "Deep Generative Models for Materials Discovery," *Frontiers in Materials*, vol. 9, 2022. doi: 10.3389/fmats.2022.865270.

[4] Nahom Aymere Hailu, "Predicting Materials' Bulk Modulus with Machine Learning," Ph.D. dissertation, [Institution Name], [City, State/Country], 2019.

[5] Samjukta Sinha1*, Prabhat Das2,3, "Random forest regression model for predicting mechanical properties," *arXiv preprint arXiv:2511.02290*, Nov. 2025.

[6] Manjunadh Kandavalli, Abhishek Agarwal, Ansh Poonia, Modalavalasa Kishor Kameswari Prasada Rao Ayyagari , "Design of high bulk moduli high entropy alloys using machine learning," *Scientific Reports*, vol. 13, 2023. doi: 10.1038/s41598-023-47181-x.

[7] Liying Cui a,‡, Qing Li a,‡, Yanchang Zhang a, Jiao Zhang a, Zhe Wang a, Jiankang Chen a, Bing Zheng a, "Machine learning-assisted high-throughput screening of 2D materials," *PubMed Central (PMC)*, PMC11575640, 2024.

[8] Caiyuan Ye, Yuzhi Wang, Xintian Xie, Tiannian Zhu, Jiaxuan Liu, Yuqing He, Lili Zhang, Junwei Zhang, Zhong Fang, Lei Wang, Zhipan Liu, Hongming Weng Quansheng Wu , "Materials discovery acceleration using conditional generative models," *npj Computational Materials*, vol. 11, 2025. doi: 10.1038/s41524-025-01930-w.

[9] A. Jain et al., "Commentary: The Materials Project: A materials genome approach to accelerating materials innovation," *APL Materials*, vol. 1, no. 1, p. 011002, Jul. 2013.

[10] S. P. Ong et al., "Python Materials Genomics (pymatgen): A robust, open-source python library for materials analysis," *Computational Materials Science*, vol. 68, pp. 314–319, Feb. 2013.